

Spark Project

Name: Mohamed Ahmed Fathy

Business concept behind the project :

ITI want to develop the job fair process and want to target international companies for student to work abroad to increase the number of people working abroad to solve the dollar issue, they want a fact table to be made over twitter to know all the tweets that are recruiting data engineers all over the world for the next 9 months, the script need to run every 10 minute and get real time data of the last 10 minute tweets talking about data engineer .

Plan in the project:

ITI isn't interested on all tweets that has word data engineer, so we need to filter on the tweets that included Job_Alert or Hiring word beside data engineer, first made several scripts to be run in 1 script , and that script with cron job will run every 10 minute and should retrieve the tweets for last 10 minutes.

List of scripts : [twitter_pipeline.py , twitter_listener.py , create_dimension.hql , twitter_stream.py , load_tweet_data.hql , spark_job.py]

Note: this code runs every 10 minutes, streaming not always opened, and in spark I make 2 tables separate, 1 in HDFS with all data, and in hive landing table which has only the new ones that we insert in dimension, there will be another solution in the folder with opened streaming and hive table on top of hive

Let's take a look on each script :

1- twitter_pipeline.py

first to run it : `python3 twitter_pipeline.py`

that is just a script that lists all other scripts inside it to run it automatically when we put it in cron job or airflow

note: twitter_listener must run in another terminal because its blocking script that's why we use .Popen
screen of it :

```
1 import os
2 import subprocess
3
4
5 # Set up Twitter stream listener
6 twitter_listener_command = ['python3', 'twitter_listener.py']
7 twitter_listener_process = subprocess.Popen(twitter_listener_command)
8
9 #Create Landing table if not exist
10 load_tweet_data_command = ['hive', '-f', 'load_tweet_data.hql']
11 subprocess.run(load_tweet_data_command, check=True)
12
13 # Stream data into Hive table
14 twitter_stream_command = ['/opt/spark-3.1.2-bin-hadoop3.2/bin/spark-submit', 'twitter_stream.py']
15 subprocess.run(twitter_stream_command, check=True)
16
17 # Create Hive dimension tables
18 hive_dimension_command = ['hive', '-f', 'create_dimension.hql']
19 subprocess.run(hive_dimension_command, check=True)
20
21 # Run Spark job
22 spark_job_command = ['/opt/spark-3.1.2-bin-hadoop3.2/bin/spark-submit', 'spark_job.py']
23 subprocess.run(spark_job_command, check=True)
24 |
```

2- twitter_listener.py

first we run it as : `python3 twitter_listener.py`

this scripts is used to listen to twittar api using your token, to extract all the tweets in the interval that contains the word `data_engineer` , so we can filter it to get the companies hiring.

The fields extracted from this tweet is :

tweet_id : which is id of each tweet

tweet_text : which is the text of the tweet

created_at : which is the time stamp of the tweet

location_name: which is the city that tweet was produced from if determined

language: which is language of the tweet

author name : which is the name of the account who wrote tweet

author_verified : which is if the account is verified or not, as if we have many tweets , we can filter for verified as it's more trusted source

public_metrics: which is all metrics of tweet as likes,retweets,etc
we put all those attributes in a forloop and save it in json formate to parse

it in the other script to convert it as columns like :

```

for data in json_response['data']:
    tweet_id = data['id']
    tweet_text = data['text']
    created_at = data['created_at']
    if 'geo' in data:
        place_id = data['geo'].get('place_id', '')
        location_info = next((item for item in json_response['includes']['places'] if item['id'] == place_id), {})
        location_name = location_info.get('name', '').encode('utf-8')
    else:
        location_name = ''.encode('utf-8')
    language = data['lang']
    author_id = data['author_id']
    author_info = next(item for item in json_response['includes']['users'] if item['id'] == author_id)
    author_name = author_info['name'].encode('utf-8')
    author_verified = author_info['verified']
    public_metrics = data['public_metrics']

    # Encode each field in utf-8 format
    tweet_text = tweet_text.encode('utf-8')
    created_at = created_at.encode('utf-8')
    language = language.encode('utf-8')
    author_name = author_name.decode().encode('utf-8')
    location_name = location_name.decode().encode('utf-8')
    public_metrics = json.dumps(public_metrics).encode('utf-8')
    # Construct a dictionary with key-value pairs
    tweet_dict = {
        "tweet_id": tweet_id,
        "tweet_text": tweet_text.decode(),
        "created_at": created_at.decode(),
        "location_name": location_name.decode(),
        "language": language.decode(),
        "author_name": author_name.decode(),
        "author_verified": author_verified,
        "public_metrics": public_metrics.decode()
    }

    # Convert dictionary to a JSON string
    row = json.dumps(tweet_dict)

    print("Sending:", row)
    clientsocket.send(row.encode('utf-8'))

```

VIP NOTE: in this listener I wanted to make the start time = the current time – 11 minute, and end time is current time – 1, so if this script runs every 10 minute, it will get all the data of the day, and made maximum 20 to not use a lot of twittar resources, but last 10 minute may or may not have tweets, and may have so few, so for only testing purpose, made start time as current time – 2hours just for testing.

Like:

```
keyword = "Data_Engineer lang:en"
# set the end time to the current time minus 1 minute
end_time = datetime.datetime.utcnow() - datetime.timedelta(minutes=1)
end_time = end_time.isoformat(timespec='milliseconds') + 'Z'

# set the start time to the current time minus 10 minutes
start_time = datetime.datetime.utcnow() - datetime.timedelta(minutes=10)
start_time = start_time.isoformat(timespec='milliseconds') + 'Z'
max_results = 20

url = create_url(keyword, start_time, end_time, max_results)
json_response = connect_to_endpoint(url[0], headers, url[1])
```

Also didn't make it always opened because it crashes, in stead made it run for last 10 minute and the script will run every 10 minute of laptop, so like that it gets all data.

3- load_tweet_data

To run this script we use :

hive -f load_tweet_data.hql

this script is so simple, it creates the landing table in hive to load data inside it, don't worry if tables already exist, this script won't make any error because it create database and table only if it doesn't exist.

```
-- Use fathy database
USE fathy;

-- Create landing_table table if not exists
CREATE TABLE IF NOT EXISTS landing_table (
  tweet_id string,
  tweet_text string,
  created_at string,
  location string,
  language string,
  author_name string,
  author_verified string,
  retweet_count string,
  reply_count string,
  like_count string,
  quote_count string,
  tweet_type string,
  work_location string,
  year int,
  month int,
  day int,
  hour int
) STORED AS PARQUET;
```

4- twitter_stream.py

to run it : /opt/spark-3.1.2-bin-hadoop3.2/bin/spark-submit

twitter_stream.py

that is the main stream that reads data from the twitter listener and clean them to be put into parquet and hive table that we just made in last script.

First we open a spark session and set configuration as

```
import pyspark
from pyspark.sql import SparkSession
from pyspark.sql.functions import *
from pyspark.sql.types import *
import time

spark = SparkSession.builder \
    .appName("myApp") \
    .enableHiveSupport() \
    .getOrCreate()
spark.conf.set("hive.metastore.uris", "thrift://localhost:5432")
spark.conf.set("spark.sql.sources.default", "parquet")
```

Note: hive for me is location on port 5432, maybe different in other docker machine so please note to change the port that you have from this command : `sudo find / -name hive-site.xml` and then cat it

Now we will load the port, and connect to twitter_listener through it, then start the streaming and sleep a little until all data come!:

```
tweet_df = spark \
    .readStream \
    .format("socket") \
    .option("host", "127.0.0.1") \
    .option("port", 7777) \
    .load()
tweet_df_string = tweet_df.selectExpr("CAST(value AS STRING)")
writeTweet = tweet_df_string.writeStream \
    .outputMode("append") \
    .format("memory") \
    .queryName("tweetquery") \
    .trigger(processingTime='2 seconds') \
    .start()
time.sleep(60)
```

Now made few definitions as we will need it later like the schema of data that will come, function to extract data and put it in another columns like filtering for job_alert tweets and put it in a column

```
schema = StructType([
    StructField("tweet_id", StringType()),
    StructField("tweet_text", StringType()),
    StructField("created_at", StringType()),
    StructField("location", StringType()),
    StructField("language", StringType()),
    StructField("author_name", StringType()),
    StructField("author_verified", StringType()),
    StructField("public_metrics", StringType())
])

def extract_tweet_type(text):
    if "hiring" in text.lower() or "job alert" or "job_alert" in text.lower():
        return "Job Alert"
    else:
        return "Other"

def extract_work_location(text):
    if "remote" in text.lower():
        return "Remote"
    elif "work from home" in text.lower():
        return "Work From Home"
    elif "premise" in text.lower() or "office" in text.lower():
        return "On Office"
    else:
        return "Other"
```


Now we made all transformation we need in data , split the coming data into rows, every row has the columns , then read it into an rdd .

Then :split the created at into 4 columns as year month day hour , also the metrics split everyone into a new column , and applied the functions we just made to search for data in tweet text and insert some certain values.

```
df = spark.sql("select * from tweetquery")

rdd = df.rdd.flatMap(lambda row: row.value.replace('{', '###').split('###'))
df = spark.read.json(rdd, schema=schema)
df = df.withColumn("public_metrics", split(df.public_metrics, ",") \
    .selectExpr("public_metrics[0] as retweet_count", "public_metrics[1] as reply_count", "public_metrics[2] as like_count",
"public_metrics[3] as quote_count") \
    .drop("public_metrics"))
df = df.withColumn("retweet_count", regexp_extract(regexp_replace("retweet_count", "[^0-9]", ""), "\d+", 0).cast("integer"))
df = df.withColumn("reply_count", regexp_extract(regexp_replace("reply_count", "[^0-9]", ""), "\d+", 0).cast("integer"))
df = df.withColumn("like_count", regexp_extract(regexp_replace("like_count", "[^0-9]", ""), "\d+", 0).cast("integer"))
df = df.withColumn("quote_count", regexp_extract(regexp_replace("quote_count", "[^0-9]", ""), "\d+", 0).cast("integer"))
df = df.withColumn("tweet_type", tweet_type_udf(df.tweet_text)).withColumn("work_location", work_location_udf(df.tweet_text))
df = df.withColumn("location", when(df.location.isNull(), "UnDetermined").otherwise(df.location))
df = df.withColumn("year", year("created_at")) \
    .withColumn("month", month("created_at")) \
    .withColumn("day", dayofmonth("created_at")) \
    .withColumn("hour", hour("created_at"))
```

now we can save it into parquet and load it into hive table we just made:

```
df.write.mode("append").partitionBy("year", "month", "day", "hour").parquet("/twitter-landing-data/extracted_tweet_data.parquet")
df.write.insertInto("fathy.landing_table")
```

and that is the schema of final landing table :

```
root
|-- tweet_id: string (nullable = true)
|-- tweet_text: string (nullable = true)
|-- created_at: string (nullable = true)
|-- location: string (nullable = true)
|-- language: string (nullable = true)
|-- author_name: string (nullable = true)
|-- author_verified: string (nullable = true)
|-- retweet_count: integer (nullable = true)
|-- reply_count: integer (nullable = true)
|-- like_count: integer (nullable = true)
|-- quote_count: integer (nullable = true)
|-- tweet_type: string (nullable = true)
|-- work_location: string (nullable = true)
|-- year: integer (nullable = true)
|-- month: integer (nullable = true)
|-- day: integer (nullable = true)
|-- hour: integer (nullable = true)
```

5- create_dimension.hql

this will create the dimensions it self, this will create the database if doesn't exist called dimension, then insert the required data into the dimension.

Made 4 dimensions:

a) tweet_text_raw : this table is for the text it self of the tweet as:

```
CREATE TABLE IF NOT EXISTS tweet_text_raw (
    tweet_id STRING,
    tweet_text STRING
)
PARTITIONED BY (year INT, month INT, day INT, hour INT);

set hive.exec.dynamic.partition.mode=nonstrict;

INSERT INTO TABLE tweet_text_raw PARTITION (year, month, day, hour)
SELECT tweet_id, tweet_text, year, month, day, hour
FROM fathy.landing_table;
```

Data in tweet_text_raw :

```
hive> select * from tweet_text_raw;
OK
1653351788085657600 RT @huangyun_122: that means a lot!! 2023 5 2 10
NULL 2023 5 2 10
I am a data engineer , covering from ETL, data modelling , even Reporting as well NULL 2023 5 2 10
NULL 2023 5 2 10
if all of the who... NULL 2023 5 2 10
1653351393821069312 Big Data Engineer Salary And Job Description 2023 - Gitnux Blog: Gain expertise in big data tools and
SQL databases like MongoDB and Cassandra. These ... #bigdata #cto https://t.co/0wCbnfKJFL 2023 5 2 10
1653350994527555584 We are Hiring Lead Data Engineer...!! 2023 5 2 10
. NULL 2023 5 2 10
. NULL 2023 5 2 10
. NULL 2023 5 2 10
#hiring #sql #dwh #sqlperformance #azuredatabricks #azuredatalakestorage #ssis https://t.co/SEpwGy6gLG
1653348569624072218 Remote Sr Data Engineer (Java/ Python/ React)- Marketing Job at TrueLogic -Jobsclub https://t.co/0JMP
1653347662370947073 We have great #jobopportunity for Azure Data Engineer with our MNC client 2023 5 2
NULL 2023 5 2 10
#azuredataengineer #pyspark #python #spark #sql #helpjobseekers #helpingeachother #jobhelpgroup #jobsearch #opportunitiesfor
.co/712hLBusTD NULL 2023 5 2 10
1653346100579844097 EasyPost is hiring Senior Data Analytics Engineer 2023 5 2 10
NULL 2023 5 2 10
@remote #wfh NULL 2023 5 2 10
python #api #gcp #sql NULL 2023 5 2 10
NULL 2023 5 2 10
#tech #softwareengineer #jobs NULL 2023 5 2 10
https://t.co/SWAJz6JHR9 NULL 2023 5 2 10
1653351788085657600 RT @huangyun_122: that means a lot!! 2023 5 2 10
NULL 2023 5 2 10
I am a data engineer , covering from ETL, data modelling , even Reporting as well NULL 2023 5 2 10
NULL 2023 5 2 10
if all of the who... NULL 2023 5 2 10
1653351393821069312 Big Data Engineer Salary And Job Description 2023 - Gitnux Blog: Gain expertise in big data tools and
SQL databases like MongoDB and Cassandra. These ... #bigdata #cto https://t.co/0wCbnfKJFL 2023 5 2 10
1653350994527555584 We are Hiring Lead Data Engineer...!! 2023 5 2 10
. NULL 2023 5 2 10
. NULL 2023 5 2 10
. NULL 2023 5 2 10
#hiring #sql #dwh #sqlperformance #azuredatabricks #azuredatalakestorage #ssis https://t.co/SEpwGy6gLG
1653348569624072218 Remote Sr Data Engineer (Java/ Python/ React)- Marketing Job at TrueLogic -Jobsclub https://t.co/0JMP
1653347662370947073 We have great #jobopportunity for Azure Data Engineer with our MNC client 2023 5 2
NULL 2023 5 2 10
#azuredataengineer #pyspark #python #spark #sql #helpjobseekers #helpingeachother #jobhelpgroup #jobsearch #opportunitiesfor
.co/712hLBusTD NULL 2023 5 2 10
1653346100579844097 EasyPost is hiring Senior Data Analytics Engineer 2023 5 2 10
NULL 2023 5 2 10
@remote #wfh NULL 2023 5 2 10
python #api #gcp #sql NULL 2023 5 2 10
```

b) author_raw : this dimension will have the data for author as:

```
CREATE TABLE IF NOT EXISTS author_raw (  
    tweet_id STRING,  
    author_name STRING,  
    author_verified STRING,  
    location STRING  
)  
PARTITIONED BY (year INT, month INT, day INT, hour INT);  
  
INSERT INTO author_raw PARTITION (year, month, day, hour)  
SELECT tweet_id,  
author_name, author_verified, location, year, month, day, hour  
FROM fathy.landing_table;
```

Data in author_raw :

```
hive> select * from author_raw;  
OK  
1653351788085657600    Zihao Wang    false    UnDetermined    2023    5    2    10  
1653351393821069312    Suriya Subramanian    false    UnDetermined    2023    5    2    10  
1653350994527555584    TechoMinds    false    UnDetermined    2023    5    2    10  
1653348569624072218    Name Book    false    UnDetermined    2023    5    2    10  
1653347662370947073    Rashmi Raj    false    UnDetermined    2023    5    2    10  
1653346100579844097    EchoJobs | Software Engineer Jobs    false    UnDetermined    2023    5    2    10  
1653351788085657600    Zihao Wang    false    UnDetermined    2023    5    2    10  
1653351393821069312    Suriya Subramanian    false    UnDetermined    2023    5    2    10  
1653350994527555584    TechoMinds    false    UnDetermined    2023    5    2    10  
1653348569624072218    Name Book    false    UnDetermined    2023    5    2    10  
1653347662370947073    Rashmi Raj    false    UnDetermined    2023    5    2    10  
1653346100579844097    EchoJobs | Software Engineer Jobs    false    UnDetermined    2023    5    2    10  
1653361351878750208    Debo Dash    false    UnDetermined    2023    5    2    11  
1653353762784722944    wamani jacob    false    UnDetermined    2023    5    2    11  
1653353641930022913    Edureka false    UnDetermined    2023    5    2    11  
1653365984164728833    TMJ-CZE Jobs    false    UnDetermined    2023    5    2    11  
1653364490832936961    Jobs Finder    false    UnDetermined    2023    5    2    11  
1653361979199995904    Sankhyana Consultancy Services-Kenya    false    UnDetermined    2023    5    2    11  
1653361351878750208    Debo Dash    false    UnDetermined    2023    5    2    11  
1653353762784722944    wamani jacob    false    UnDetermined    2023    5    2    11  
1653353641930022913    Edureka false    UnDetermined    2023    5    2    11  
165335107737124865    jessica👩    false    UnDetermined    2023    5    2    9  
1653332761271484422    bun.bun.👶    false    UnDetermined    2023    5    2    9  
1653332630484713472    wen👩👶 false    UnDetermined    2023    5    2    9  
1653332288988487684    Ricardo Vázquez    false    UnDetermined    2023    5    2    9  
1655161185694633985    Bazzer360    false    UnDetermined    2023    5    7    10  
1655160842562818048    Tyler F👶    false    UnDetermined    2023    5    7    10  
1655158034962829313    EchoJobs | Software Engineer Jobs    false    UnDetermined    2023    5    7    10  
1655178164652351489    EchoJobs | Software Engineer Jobs    false    UnDetermined    2023    5    7    11  
1655184810875269120    Your Agile Team    true    UnDetermined    2023    5    7    12  
Time taken: 2.304 seconds, Fetched: 30 row(s)  
hive> describe author_raw;  
OK  
tweet_id                string  
author_name              string  
author_verified          string  
location                 string  
year                     int  
month                    int  
day                      int  
hour                     int  
  
# Partition Information  
# col_name               data_type      comment  
year                     int
```

c) tweet_properties_raw: this will have some information about tweet

```
CREATE TABLE IF NOT EXISTS tweet_properties_raw (  
    tweet_id STRING,  
    work_location STRING,  
    tweet_type STRING,  
    created_at STRING,  
    language STRING  
)  
PARTITIONED BY (year INT, month INT, day INT, hour INT);  
  
INSERT INTO tweet_properties_raw PARTITION (year, month, day, hour)  
SELECT tweet_id,  
    work_location,  
    tweet_type,  
    created_at,  
    language,  
    year,  
    month,  
    day,  
    hour  
FROM fathy.landing_table;
```

Data in tweet_properties_raw:

```
hive> select * from tweet_properties_raw;  
OK
```

1653351788085657600	Other	Job Alert	2023-05-02T10:52:38.000Z	en	2023	5	2	10
1653351393821069312	Other	Job Alert	2023-05-02T10:51:04.000Z	en	2023	5	2	10
1653350994527555584	Other	Job Alert	2023-05-02T10:49:29.000Z	en	2023	5	2	10
1653348569624072218	Remote	Job Alert	2023-05-02T10:39:51.000Z	en	2023	5	2	10
1653347662370947073	Other	Job Alert	2023-05-02T10:36:15.000Z	en	2023	5	2	10
1653346100579844097	Remote	Job Alert	2023-05-02T10:30:02.000Z	en	2023	5	2	10
1653351788085657600	Other	Job Alert	2023-05-02T10:52:38.000Z	en	2023	5	2	10
1653351393821069312	Other	Job Alert	2023-05-02T10:51:04.000Z	en	2023	5	2	10
1653350994527555584	Other	Job Alert	2023-05-02T10:49:29.000Z	en	2023	5	2	10
1653348569624072218	Remote	Job Alert	2023-05-02T10:39:51.000Z	en	2023	5	2	10
1653347662370947073	Other	Job Alert	2023-05-02T10:36:15.000Z	en	2023	5	2	10
1653346100579844097	Remote	Job Alert	2023-05-02T10:30:02.000Z	en	2023	5	2	10
1653361351878750208	Other	Job Alert	2023-05-02T11:30:38.000Z	en	2023	5	2	11
1653353762784722944	Other	Job Alert	2023-05-02T11:00:29.000Z	en	2023	5	2	11
1653353641930022913	Other	Job Alert	2023-05-02T11:00:00.000Z	en	2023	5	2	11
1653365984164728833	Other	Job Alert	2023-05-02T11:49:03.000Z	en	2023	5	2	11
1653364490832936961	Other	Job Alert	2023-05-02T11:43:07.000Z	en	2023	5	2	11
1653361979199995904	Other	Job Alert	2023-05-02T11:33:08.000Z	en	2023	5	2	11
1653361351878750208	Other	Job Alert	2023-05-02T11:30:38.000Z	en	2023	5	2	11
1653353762784722944	Other	Job Alert	2023-05-02T11:00:29.000Z	en	2023	5	2	11
1653353641930022913	Other	Job Alert	2023-05-02T11:00:00.000Z	en	2023	5	2	11
165335107737124865	Other	Job Alert	2023-05-02T09:46:21.000Z	en	2023	5	2	9
165333276127148442	Other	Job Alert	2023-05-02T09:37:02.000Z	en	2023	5	2	9
1653332630484713472	Other	Job Alert	2023-05-02T09:36:31.000Z	en	2023	5	2	9
1653332288988487684	Other	Job Alert	2023-05-02T09:35:09.000Z	en	2023	5	2	9
1655161185694633985	Other	Job Alert	2023-05-07T10:42:32.000Z	en	2023	5	7	10
1655160842562818048	Other	Job Alert	2023-05-07T10:41:10.000Z	en	2023	5	7	10
1655158034962829313	Remote	Job Alert	2023-05-07T10:30:01.000Z	en	2023	5	7	10
1655178164652351489	Other	Job Alert	2023-05-07T11:50:00.000Z	en	2023	5	7	11
1655184810875269120	Other	Job Alert	2023-05-07T12:16:25.000Z	en	2023	5	7	12

d) metrics_raw : which has data inside metrics_raw

```
CREATE TABLE IF NOT EXISTS metrics_raw (
  tweet_id STRING,
  retweet_count INT,
  reply_count INT,
  like_count INT,
  quote_count INT
)
PARTITIONED BY (year INT, month INT, day INT, hour INT);

INSERT INTO metrics_raw PARTITION (year, month, day, hour)
SELECT tweet_id,
  retweet_count,
  reply_count,
  like_count,
  quote_count,
  year,
  month ,
  day ,
  hour
FROM fathy.landing_table;
```

Data in metrics_raw:

```
hive> select * from metrics_raw;
OK
1653351788085657600      1      0      0      0      2023      5      2      10
1653351393821069312      0      0      0      0      2023      5      2      10
1653350994527555584      0      0      1      0      2023      5      2      10
1653348569624072218      0      0      0      0      2023      5      2      10
1653347662370947073      0      0      0      0      2023      5      2      10
1653346100579844097      0      0      0      0      2023      5      2      10
1653351788085657600      1      0      0      0      2023      5      2      10
1653351393821069312      0      0      0      0      2023      5      2      10
1653350994527555584      0      0      1      0      2023      5      2      10
1653348569624072218      0      0      0      0      2023      5      2      10
1653347662370947073      0      0      0      0      2023      5      2      10
1653346100579844097      0      0      0      0      2023      5      2      10
1653361351878750208      1      0      0      0      2023      5      2      11
1653353762784722944      0      0      2      0      2023      5      2      11
1653353641930022913      0      0      0      0      2023      5      2      11
1653365984164728833      0      0      0      0      2023      5      2      11
1653364490832936961      0      0      0      0      2023      5      2      11
1653361979199995904      0      0      1      0      2023      5      2      11
1653361351878750208      1      0      0      0      2023      5      2      11
1653353762784722944      0      0      2      0      2023      5      2      11
1653353641930022913      0      0      0      0      2023      5      2      11
165335107737124865      16      0      0      0      2023      5      2      9
1653332761271484422      16      0      0      0      2023      5      2      9
1653332630484713472      16      0      0      0      2023      5      2      9
1653332288988487684      1      0      0      0      2023      5      2      9
1655161185694633985      0      1      2      0      2023      5      7      10
1655160842562818048      0      1      2      0      2023      5      7      10
1655158034962829313      0      0      0      0      2023      5      7      10
1655178164652351489      0      0      0      0      2023      5      7      11
1655184810875269120      0      0      0      0      2023      5      7      12
Time taken: 0.235 seconds, Fetched: 30 row(s)
hive> describe metrics_raw;
OK
tweet_id                string
retweet_count           int
reply_count             int
like_count              int
quote_count             int
year                    int
month                   int
day                     int
hour                    int
```

Now all dimensions are ready to make the fact table!

6- spark_job

to run it: `/opt/spark-3.1.2-bin-hadoop3.2/bin/spark-submit spark_job.py`

this will read the data from hive's dimension and make fact tables

first read files from hive as :

```
df_tweet_text = spark.sql("SELECT tweet_id, tweet_text FROM dimension.tweet_text_raw WHERE tweet_text IS NOT NULL")
df_tweet_properties = spark.sql("SELECT * FROM dimension.tweet_properties_raw")
df_tweet_author = spark.sql("SELECT * FROM dimension.author_raw")
df_tweet_metrics = spark.sql("SELECT * FROM dimension.metrics_raw")
```

now we will make fact tables from this hive as:

```
df_tweet_properties_filtered = df_tweet_properties.filter(lower(df_tweet_properties.tweet_type) == 'job alert')
df_tweet_metrics_properties = df_tweet_metrics.select('tweet_id', 'retweet_count', 'like_count')
df_tweet_metrics_properties = df_tweet_properties_filtered.join(df_tweet_metrics_properties, 'tweet_id', 'inner')
df_tweet_text_filtered = df_tweet_text.select('tweet_id', 'tweet_text')
df_tweet_engagement = df_tweet_metrics_properties.join(df_tweet_text_filtered, 'tweet_id', 'inner')
df_tweet_engagement_hour = df_tweet_engagement.groupBy('tweet_type', 'language', 'work_location', 'hour', 'day', 'month', 'year').agg({'like_count':
'sum', 'retweet_count': 'sum'})

df_tweet_engagement_hour = df_tweet_engagement_hour.withColumnRenamed('sum(like_count)', 'like_count').withColumnRenamed('sum(retweet_count)',
'retweet_count')

spark.sql("CREATE DATABASE IF NOT EXISTS twitter_processed_data")
df_tweet_engagement_hour.write.mode('overwrite').saveAsTable('twitter_processed_data.tweet_engagement_hour_processed')
df_tweet_engagement.write.mode('overwrite').saveAsTable('twitter_processed_data.tweet_engagement_processed')
```

then made another fact table as :

```
df_tweet_location_filtered = df_tweet_properties.filter(lower(df_tweet_properties.tweet_type) == 'job alert')
df_tweet_metrics_properties = df_tweet_metrics.select('tweet_id', 'retweet_count', 'like_count', 'reply_count')
df_tweet_metrics_properties = df_tweet_location_filtered.join(df_tweet_metrics_properties, 'tweet_id', 'inner')
df_tweet_author_filtered = df_tweet_author.select('tweet_id', 'location')
df_tweet_author_filtered = df_tweet_author_filtered.join(df_tweet_metrics_properties, 'tweet_id', 'inner')

df_tweet_location_engagement = df_tweet_author_filtered.groupBy('location', 'hour', 'day', 'month', 'year').agg({
    'retweet_count': 'sum',
    'like_count': 'sum',
    'reply_count': 'sum',
    'tweet_type': 'count'
}).withColumnRenamed('count(tweet_type)', 'tweet_count').withColumnRenamed('sum(like_count)', 'likes') \
.withColumnRenamed('sum(reply_count)', 'replies').withColumnRenamed('sum(retweet_count)', 'retweets')

df_tweet_location_engagement.write.mode('overwrite').saveAsTable('twitter_processed_data.tweet_location_processed')
```

Now let's take a look on the data it self for each fact table as :

1) tweet_engagement_processed :

this fact table tweets that is only looking for jobs like id of tweet and it's properties, like language? Timestamp? When? Likes? Retweets? And so on to make queries on it filtered for job alert and here sample of data:

```
hive> describe tweet_engagement_processed;
OK
tweet_id          string
work_location     string
tweet_type        string
created_at        string
language          string
year              int
month             int
day               int
hour              int
retweet_count     int
like_count        int
tweet_text        string
Time taken: 0.029 seconds, Fetched: 12 row(s)
hive> select * from tweet_engagement_processed;
OK
1653347662370947073  Other  Job Alert  2023-05-02T10:36:15.000Z  en  2023  5  2  10  0  0  We have great #joboportur
ity for Azure Data Engineer with our MMC client
1653347662370947073  Other  Job Alert  2023-05-02T10:36:15.000Z  en  2023  5  2  10  0  0  We have great #joboportur
ity for Azure Data Engineer with our MMC client
1653347662370947073  Other  Job Alert  2023-05-02T10:36:15.000Z  en  2023  5  2  10  0  0  We have great #joboportur
ity for Azure Data Engineer with our MMC client
1653347662370947073  Other  Job Alert  2023-05-02T10:36:15.000Z  en  2023  5  2  10  0  0  We have great #joboportur
```

2) tweet_engagement_hour_processed

Same table as upove but grouped by hour/day/month/year to get more aggregated data

```
hive> select * from tweet_engagement_hour_processed;
OK
Job Alert  en  Other  11  7  5  2023  0  0
Job Alert  en  Other  9  2  5  2023  0  49
Job Alert  en  Other  11  2  5  2023  17  8
Job Alert  en  Other  10  2  5  2023  8  8
Job Alert  en  Remote  10  2  5  2023  0  0
Job Alert  en  Remote  10  7  5  2023  0  0
Job Alert  en  Other  12  7  5  2023  0  0
Job Alert  en  Other  10  7  5  2023  4  0
Time taken: 0.121 seconds, Fetched: 8 row(s)
hive> describe tweet_engagement_hour_processed;
OK
tweet_type        string
language          string
work_location     string
hour              int
day               int
month             int
year              int
like_count        bigint
retweet_count     bigint
Time taken: 0.041 seconds, Fetched: 9 row(s)
```

3) now 3rd fact table which is tweet_egagement_location_processed;

which groups for every country the tweets, and if it's Undetermined it says undetermined means twitter couldn't fetch the country of the person who made the tweet.

As my data set here is for only last 10minute so I don't have much data, so I got only undetermined, but if a lot of data should be different countries as twittar doesn't allow giving location to developers unless they allow it .

```
hive> select * from tweet_location_processed
> ;
OK
Undetermined    9      2      5      2023    0      0      49      4
Undetermined    11     2      5      2023    17     0      8      27
Undetermined    11     7      5      2023    0      0      0      1
Undetermined    12     7      5      2023    0      0      0      1
Undetermined    10     7      5      2023    4      2      0      3
Undetermined    10     2      5      2023    8      0      8      48
Time taken: 0.109 seconds, Fetched: 6 row(s)
hive> describe tweet_location_processed;
OK
location          string
hour              int
day               int
month             int
year              int
likes             bigint
replies          bigint
retweets         bigint
tweet_count      bigint
Time taken: 0.037 seconds, Fetched: 9 row(s)
hive> |
```

Now will make cron job to run this script every 10 minutes:

Now write command:

a- **crontab -e**

then write that line inside it :

***/10 * * * * python3 /home/itversity/itversity-**

material/project/twitter_pipeline.py >/dev/null 2>&1

This will run all the pipeline every 10 minute to extract data in last 10 minute